

Databricks Community Edition - Full Practice Set (30 Exercises)

Instructions:

Run each code block in Databricks Community Edition.

Observe the results, compare with expected outputs.

Exercise 1: Load a CSV file into a DataFrame

Code:

```
df = spark.read.option("header", True).option("inferSchema", True).csv("/databricks-datasets/airlines/part-00000")  
  
df.show(5)
```

Expected Result:

Shows first 5 rows of the Airlines dataset.

Exercise 2: Cache a DataFrame for faster access

Code:

```
df.cache()  
  
df.count()
```

Expected Result:

Caches the DataFrame into memory.

Exercise 3: Select specific columns from a DataFrame

Code:

```
df.select("Year", "Month", "DepDelay").show(5)
```

Expected Result:

Shows selected columns.

Exercise 4: Filter rows where DepDelay > 30 minutes

Code:

```
df.filter(df.DepDelay > 30).show(5)
```

Expected Result:

Shows flights delayed more than 30 minutes.

Exercise 5: Add a new column 'DelayCategory'

Code:

```
from pyspark.sql.functions import when
```

```
df2 = df.withColumn("DelayCategory", when(df.DepDelay > 30, "Late").otherwise("OnTime"))
```

```
df2.select("DepDelay", "DelayCategory").show(5)
```

Expected Result:

Adds column categorizing flights as Late or OnTime.

Exercise 6: Group by Origin and count flights

Code:

```
df.groupBy("Origin").count().orderBy("count", ascending=False).show(5)
```

Expected Result:

Shows airports with most flights.

Exercise 7: Sort flights by Departure Delay descending

Code:

```
df.orderBy(df.DepDelay.desc()).show(5)
```

Expected Result:

Flights with highest departure delay first.

Exercise 8: Rename column DepDelay to DepartureDelayMinutes

Code:

```
df2 = df.withColumnRenamed("DepDelay", "DepartureDelayMinutes")
```

```
df2.show(5)
```

Expected Result:

Column renamed.

Exercise 9: Drop rows with null DepDelay

Code:

```
df_clean = df.na.drop(subset=["DepDelay"])
```

```
df_clean.show(5)
```

Expected Result:

Null DepDelay rows removed.

Exercise 10: Get distinct Origin airports

Code:

```
df.select("Origin").distinct().show(5)
```

Expected Result:

Lists distinct origin airports.

Exercise 11: Save DataFrame as Delta Table

Code:

```
df.write.format("delta").mode("overwrite").save("/tmp/airlines_delta")
```

Expected Result:

Saves as Delta format.

Exercise 12: Load Delta Table into DataFrame

Code:

```
df_delta = spark.read.format("delta").load("/tmp/airlines_delta")
```

```
df_delta.show(5)
```

Expected Result:

Loads Delta format back.

Exercise 13: Create a SQL table from Delta location

Code:

```
spark.sql("CREATE TABLE IF NOT EXISTS airlines_delta USING DELTA LOCATION  
'/tmp/airlines_delta'")
```

Expected Result:

Creates SQL table linked to Delta files.

Exercise 14: Update records in Delta Table

Code:

```
spark.sql("UPDATE airlines_delta SET DepDelay = 0 WHERE DepDelay IS NULL")
```

Expected Result:

Updates NULL DepDelay values.

Exercise 15: Time Travel to old Delta Table version

Code:

```
df_old = spark.read.format("delta").option("versionAsOf", 0).load("/tmp/airlines_delta")  
df_old.show(5)
```

Expected Result:

Reads earlier version.

Exercise 16: Merge new data into Delta Table

Code:

```
from delta.tables import DeltaTable
```

```
deltaTable = DeltaTable.forPath(spark, "/tmp/airlines_delta")
```

```
deltaTable.alias("old").merge(
```

```
df.alias("new"), "old.FlightNum = new.FlightNum"  
).whenMatchedUpdateAll().whenNotMatchedInsertAll().execute()
```

Expected Result:

Upserts data into Delta Table.

Exercise 17: Optimize Delta Table

Code:

```
spark.sql("OPTIMIZE airlines_delta")
```

Expected Result:

Improves Delta query speed.

Exercise 18: Vacuum old files from Delta Table

Code:

```
spark.sql("VACUUM airlines_delta RETAIN 168 HOURS")
```

Expected Result:

Cleans up old files safely.

Exercise 19: Add a new column to Delta Table

Code:

```
spark.sql("ALTER TABLE airlines_delta ADD COLUMNS (NewColumn STRING)")
```

Expected Result:

New column added.

Exercise 20: Drop a column from Delta Table

Code:

```
spark.sql("ALTER TABLE airlines_delta DROP COLUMN NewColumn")
```

Expected Result:

Column removed from Delta Table.

Exercise 21: Run SQL query on Delta Table

Code:

```
spark.sql("SELECT Carrier, AVG(DepDelay) AS AvgDepDelay FROM airlines_delta GROUP BY Carrier ORDER BY AvgDepDelay DESC").show()
```

Expected Result:

Average delays per Carrier.

Exercise 22: Create managed Delta Table from DataFrame

Code:

```
df.write.saveAsTable("managed_airlines_table")
```

Expected Result:

Creates managed SQL table.

Exercise 23: Drop SQL table

Code:

```
spark.sql("DROP TABLE IF EXISTS managed_airlines_table")
```

Expected Result:

Deletes SQL table.

Exercise 24: Create a temporary SQL View

Code:

```
df.createOrReplaceTempView("temp_view_flights")
```

```
spark.sql("SELECT * FROM temp_view_flights LIMIT 5").show()
```

Expected Result:

Creates view for easy querying.

Exercise 25: Check Delta Table history

Code:

```
spark.sql("DESCRIBE HISTORY airlines_delta").show(truncate=False)
```

Expected Result:

Shows table modification history.

Exercise 26: Top 5 flights with highest Arrival Delay

Code:

```
df.orderBy(df.ArrDelay.desc()).select("FlightNum", "ArrDelay").show(5)
```

Expected Result:

Top delayed flights displayed.

Exercise 27: Find flights with negative DepDelay (early departures)

Code:

```
df.filter(df.DepDelay < 0).select("FlightNum", "DepDelay").show(5)
```

Expected Result:

Early departing flights listed.

Exercise 28: Group by Carrier and calculate avg DepDelay

Code:

```
df.groupBy("Carrier").avg("DepDelay").orderBy("avg(DepDelay)").show(5)
```

Expected Result:

Best and worst airlines by delay.

Exercise 29: Register DataFrame as SQL Table then query

Code:

```
df.write.format("delta").mode("overwrite").saveAsTable("airport_summary")
```

```
spark.sql("SELECT Origin, COUNT(*) FROM airport_summary GROUP BY Origin ORDER BY  
COUNT(*) DESC").show(5)
```

Expected Result:

Top airports displayed.

Exercise 30: Perform Delta Time Travel using SQL

Code:

```
spark.sql("SELECT * FROM airlines_delta VERSION AS OF 0 LIMIT 5").show()
```

Expected Result:

Reads old snapshot of table.